

Pre-Lab: Installing pavo and Friends

Complete BEFORE Wednesday May 6 — Lab 5: Pavo Introduction

EEB 187 — Ecology & Evolution of Color

Table of contents

Why this matters	1
What you need to install	2
Step 1: Install R (if you don't have it)	3
Step 2: Install RStudio	3
Step 3: Install the R packages	3
Step 4: Smoke test	4
Step 5 (optional but recommended): Create a lab project folder	5
Step 6 (RECOMMENDED): Install the segmentation tools	5
macOS	5
Windows	7
Linux	7
Recommended alternative: install uv instead of pip-installing rembg	7
If install fails entirely	9
Step 7 (REQUIRED): Identify three fish from your assigned lineage	9
Step 7a: Where to put your images	9
Step 7b: Image naming convention	10
Step 7c (RECOMMENDED): Segment your fish images BEFORE lab	10
Common errors and fixes	11
What we'll do in lab	12
Need help?	12

Why this matters

Wednesday's lab is a hands-on R session. **You must arrive with a working R + RStudio + pavo installation.** If you try to install during class, you will fall behind and miss the live walkthrough. Allocate about **30 minutes** the day before lab to do this. Most of that time is downloads.

If you hit problems, post in the EEB 187 channel — there's almost certainly a classmate fighting the same error.

💡 Don't want to install R, ImageMagick, and rembg natively? Use Docker — skip Steps 1–6.

If you already use Docker (or are happy to install [Docker Desktop](#) once), the entire lab environment is pre-baked: R 4.4 + RStudio Server + pavo + tidyverse + magick + ImageMagick + uv/rembg. **You can skip Steps 1–6 below entirely.** Jump to [Step 7](#) once your container is running.

```
git clone https://github.com/AlfaroLab/eeb187-pavo-lab.git
cd eeb187-pavo-lab
docker build -t pavo-lab . # ~10 min the first time, instant after
docker run --rm -p 8787:8787 \
  -e PASSWORD=eeb187 \
  -v "$(pwd):/home/rstudio/lab" \
  pavo-lab
```

Open <http://localhost:8787> in your browser. Log in as **rstudio** / **eeb187**. The lab repo is mounted at `/home/rstudio/lab`, so any files you save there persist on your laptop after the container exits.

To verify the pipeline (same 4/4 **hard asserts passed** smoke test as the native path), open the RStudio Terminal tab inside the container and run:

```
cd /home/rstudio/lab
Rscript scripts/verify_pipeline.R
```

Trade-offs: Docker images are large (~3 GB on disk after the build). Apple Silicon Macs and recent Windows machines run this fine. If your laptop is older or short on disk, prefer one of the cloud paths (Binder/Codespaces) instead.

What you need to install

Tool	Version	Why
R	4.3	The language pavo runs on
RStudio	Recent (2024+)	The editor we'll use
pavo	2.9	Color pattern analysis
tidyverse	Recent	Data wrangling + ggplot
magick	Recent	Image I/O backend pavo uses

Step 1: Install R (if you don't have it)

- **macOS:** <https://cran.r-project.org/bin/macosx/> — get the `.pkg` for your chip (Apple Silicon = “arm64”, Intel = “x86_64”)
- **Windows:** <https://cran.r-project.org/bin/windows/base/> — get the `.exe`
- **Linux:** use your distro's package manager (`apt install r-base`, etc.)

Verify by opening a terminal and running:

```
R --version
```

You should see something like `R version 4.3.x` or higher.

Step 2: Install RStudio

<https://posit.co/download/rstudio-desktop/> — pick the free desktop version. Install. Open it once to confirm it launches.

Step 3: Install the R packages

Open RStudio. In the Console pane (bottom-left), paste and run:

```
install.packages(c("pavo", "tidyverse", "magick", "knitr"))
```

This will take **5–15 minutes** depending on your connection. You'll see a lot of compilation output. That's normal. Wait for the prompt (`>`) to come back.

 Warning

On macOS: if you see errors about Cellar, gfortran, or clang, you need Apple's command-line tools first. Run in Terminal: `xcode-select --install`, then re-run the R install.

On Windows: if pavo fails to install, you may need to install Rtools first: <https://cran.r-project.org/bin/windows/Rtools/>

Step 4: Smoke test

In the RStudio Console, paste and run **all** of this:

```
library(pavo)
library(magick)
library(tidyverse)

# Load a test image bundled with pavo
data(sicalis)

# Run a quick check of the pavo version
packageVersion("pavo")

cat("\nIf you see a version number above ( 2.9.0) and no errors,\n")
cat("you are ready for Wednesday's lab.\n")
```

You should see:

```
[1] '2.9.0'    # or higher
```

If you see a version number above (2.9.0) and no errors,
you are ready for Wednesday's lab.

If you got that, **stop**. You're done.

Step 5 (optional but recommended): Create a lab project folder

Pick a place on your computer where you keep coursework — Documents, Desktop, wherever. Make a folder called `EEB187-pavo-lab`. On Wednesday I'll have you download a ZIP from BruinLearn that you'll unpack into this folder. Doing this in advance saves 90 seconds of fumbling on the day.

```
Documents/  
  EEB187-pavo-lab/      ← create this empty folder now
```

Step 6 (RECOMMENDED): Install the segmentation tools

Your fish images need their backgrounds removed before pavo can analyze them cleanly. The lab uses two command-line tools, both invoked from R via a one-line helper. **Install both before lab if you can** — installing during lab will eat your work time.

macOS

i Don't have Homebrew yet? Most students won't.

Homebrew (`brew`) is the macOS package manager — a one-time install that makes installing things like ImageMagick a one-liner forever after. **You only need to do this once on your laptop.**

First, install Xcode Command Line Tools

Homebrew depends on Apple's Command Line Tools (the C/C++ compiler and friends). If you've never installed them, **do this first**. In Terminal (**Applications → Utilities → Terminal**):

```
xcode-select --install
```

A small dialog box will pop up. Click **Install** (not "Get Xcode" — you don't need the full app). It downloads ~3 GB and takes 10–20 minutes depending on your internet. When it finishes, run the same command again — if it says "*command line tools are already installed*," you're done with this step.

If you've used `git`, compiled anything, or installed Xcode before, you may already have the Command Line Tools — running the command above is harmless either way.

Then install Homebrew — pick whichever option you prefer

Option A — the official installer script (most common, most documentation references this):

1. Paste this in Terminal and press return:

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

2. It will ask for your **macOS password** (the one you use to log into your Mac). Type it — characters won't show as you type, that's normal — and press return.
3. The installer downloads and configures Homebrew. Watch for a message at the very end labeled **"Next steps"**. On Apple-Silicon Macs (M1/M2/M3/M4) it will tell you to run two more commands that look like this:

```
echo 'eval "$(/opt/homebrew/bin/brew shellenv)"' >> ~/.zprofile
eval "$(/opt/homebrew/bin/brew shellenv)"
```

Run those exact commands (copy from your Terminal output, not from this guide — your username will be different). They put **brew** on your **PATH** so future Terminal sessions can find it.

Option B — the signed .pkg installer (point-and-click; nice if you don't want to paste shell commands):

1. Open <https://github.com/Homebrew/brew/releases/latest> in your browser.
2. Under **Assets**, download the file ending in **.pkg** (e.g. **Homebrew-4.x.x.pkg**).
3. Double-click the downloaded **.pkg**. The standard macOS installer opens — click through the prompts, enter your macOS password when asked, and let it finish.
4. Open a fresh Terminal window. Apple-Silicon users still need to run the two **shellenv** lines from Option A step 3 once — the **.pkg** doesn't always do this automatically.

Verify whichever option you used

```
brew --version # should print "Homebrew 4.x.x"
```

If **brew --version** says “command not found,” your shell hasn't been told about Homebrew yet. **Quit Terminal and reopen it**, then try again. If still no luck, run the two **shellenv** lines from Option A step 3.

Full official instructions, in case anything looks different on your machine: <https://brew.sh>

Once Homebrew is in place:

```
# In Terminal:
brew install imagemagick
pip install "rembg[cpu,cli]"          # or pipx install "rembg[cpu,cli]"
```

Verify both work:

```
magick -version      # should print "Version: ImageMagick 7.x..."
rembg --help         # should print rembg's usage text
```

Windows

Install [ImageMagick from the official site](#) (the “Windows Binary Release” installer), then in Command Prompt:

```
pip install "rembg[cpu,cli]"
```

Linux

```
sudo apt install imagemagick
pip install "rembg[cpu,cli]"
```

Recommended alternative: install uv instead of pip-installing rembg

Several students will have trouble with `pip install "rembg[cpu,cli]"` because rembg’s deep-learning dependencies (onnxruntime, opencv) don’t always have prebuilt wheels for the latest Python. There’s a much smoother path that **handles all of this automatically**:

```
# macOS:
brew install uv

# Linux:
curl -LsSf https://astral.sh/uv/install.sh | sh

# Windows:
powershell -ExecutionPolicy Bypass -c "irm https://astral.sh/uv/install.ps1 | iex"
```

That's it. With `uv` installed, the `segment_fish()` helper in lab will run `rembg` via `uvx --from "rembg[cpu,cli]" rembg ...` — `uv` downloads a compatible Python and `rembg` automatically the first time you call the helper, and caches both for future calls. **No more pip install errors, no Python version mismatches.**

`segment_fish()` tries `uvx` first and falls back to a direct `rembg` binary, so if you got `rembg` installed via `pip` and it's on `PATH`, that works too.

 Common roadblock 1 — Windows + pip install fails on `rembg`

If you tried `pip install "rembg[cpu,cli]"` on Windows and saw errors about Visual C++ build tools, ONNX runtime, or `opencv`: **don't fight it**. Install `uv` instead (one PowerShell command above) and skip `pip` entirely. `segment_fish()` will use `uvx` automatically.

 Common roadblock 2 — “`rembg` installed but R can't find it”

If `rembg --help` works in your terminal but R says `rembg failed (or not installed)`: R doesn't see your full `PATH`. Check what R sees:

```
Sys.getenv("PATH")
```

If `/opt/homebrew/bin` (macOS Homebrew) or `~/.local/bin` (pip user installs) isn't in there, **install `uv` and use that path instead** — `uvx` is more discoverable. Or set the `PATH` in `~/.Renviron`:

```
PATH=/opt/homebrew/bin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:~/.local/bin
```

Restart RStudio for the change to take effect.

 Common roadblock 3 — “I had to fall back to manual cropping; how does that affect my answers?”

If neither `uv` nor `rembg` work for you, the worksheet's **manual-crop path** (Step 1.5 Path 2) is fully supported. Your A values will be less dramatic than the demo (closer to 1 across the board) because some background pixels will pollute the analysis. **The directional A claim still holds** — vertical-bar fish > 1 , horizontal-stripe fish < 1 — just with smaller effect sizes. Note in your worksheet which path you used; the rubric accommodates both.

If install fails entirely

Come to lab anyway with your three lineage fish images saved locally. Use the manual-crop path. The lab is still teachable; just plan for the smaller-magnitude metrics described above.

Step 7 (REQUIRED): Identify three fish from your assigned lineage

Your group has been assigned a reef fish lineage (typically a family). Before lab on Wednesday, **pick three species from your lineage that show three visually different pattern types** — for example:

- One with strong **vertical bars** (you'll predict $A > 1$)
- One with a **horizontal stripe or band** (you'll predict $A < 1$)
- One that's **spotted, reticulate, or solid** (you'll predict $A \approx 1$)

If your lineage doesn't span all three pattern types, pick the three with the most different patterns you can find in it.

Image source — use whichever your group has been working with:

- **Fastest path** — **already done for you:** the lab bundle includes pre-curated images at `images/student-fish/<your-lineage>/` — five species per lineage, ready to use. Skip downloading and segmentation entirely if you pick three of these.
- **FishView server** — if your group has been curating images there, use that
- **FishBase** ([fishbase.se](https://www.fishbase.org/)) or **Wikimedia Commons** — search by species name; download a clear, lateral-view photo

Step 7a: Where to put your images

Use one consistent folder so all your `getimg()` calls use short, predictable paths. The convention we recommend:

OS	Recommended path
macOS	<code>~/Documents/eeb187-pavo-lab/images/student-fish/<your-lineage>/</code>
Windows	<code>Documents\eeb187-pavo-lab\images\student-fish\<your-lineage>\</code>

OS	Recommended path
Linux	<code>~/eeb187-pavo-lab/images/student-fish/<your-lineage>/</code>
Binder / Codespaces	<code>images/student-fish/<your-lineage>/</code> (use the RStudio Files-pane <i>Upload</i> button)

If you cloned the repo with `git clone https://github.com/AlfaroLab/eeb187-pavo-lab.git`, that folder *already exists* with the 5 starter species in it — just drop your own JPGs in alongside them.

💡 Why this convention?

Every code example in `walkthrough.R` and the worksheet uses paths like `images/student-fish/chaetodontidae/chaetodon-auriga.jpg`. If your folder layout matches, every code block works copy-paste with one filename change. If it doesn't, you'll be debugging file paths instead of biology.

Step 7b: Image naming convention

Use lowercase, hyphenated filenames so the helper functions and your code are robust:

- `genus-species.jpg` — e.g. `chaetodon-auriga.jpg`, `pomacanthus-imperator.jpg`
- For multiple photos of the same species: `genus-species-1.jpg`, `genus-species-2.jpg`
- **Avoid spaces, capital letters, parentheses, and special characters** — they break R file paths on some platforms

Step 7c (RECOMMENDED): Segment your fish images BEFORE lab

If you can pre-segment your 3 fish images at home, you'll spend lab time on the *biology* (predict-and-test, comparison) instead of on image processing. **Two paths, both fine:**

Path A — local with `segment_fish()` (best quality): if your install of `uv` + `imagemagick` worked above, you already have the tool. Just run in R:

```
source("scripts/walkthrough.R") # loads the segment_fish() helper
segment_fish("images/student-fish/<your-lineage>/your-fish.jpg")
# Repeat for fish 2 and 3. Output: <name>-segmented.jpg in the same folder.
```

Path B — web tool (no install needed): drag-drop your fish JPG to one of these:

- [remove.bg](#) — best known, free for low-res output (fine for this lab), no signup needed
- [photoroom.com](#) — free, no signup
- [clipdrop.co/remove-background](#) — free, no signup

The output is a transparent PNG. Save it next to the original (e.g., `chaetodon-auriga.png`). Then in R, use the `composite_to_white()` helper from `walkthrough.R` to convert the transparent PNG to a white-background JPG that pavo expects:

```
source("scripts/walkthrough.R")
composite_to_white("images/student-fish/<your-lineage>/chaetodon-auriga.png")
# Output: chaetodon-auriga-segmented.jpg in the same folder.
```

Either way, you arrive at lab with three `<species>-segmented.jpg` files ready for `getimg()`. **Don't show up empty-handed.** If you don't know your lineage assignment, ask in the channel by Tuesday evening.

Common errors and fixes

there is no package called 'pavo' — `install.packages("pavo")` didn't actually run. Re-run it; check for a typo.

unable to load shared object ... magick.so — On macOS, install ImageMagick via Homebrew first: `brew install imagemagick`, then reinstall the R magick package. On Linux: `sudo apt install libmagick++-dev`.

installation of package 'XXX' had non-zero exit status — The package failed to compile. Scroll up in the install output to see the actual error. Most common: missing system dependency. Google the error or post in the channel.

RStudio crashes on startup — Reinstall RStudio. If still failing, your R install is broken; reinstall R, then RStudio.

What we'll do in lab

You don't need to read up on pavo or do prep beyond Steps 1–7 above. Wednesday's lab walks you through everything from scratch:

1. **Demo block** — three reef fish from three families: sergeant major (vertical bars), cleaner wrasse (horizontal stripe), mandarin fish (reticulate). All three are pre-segmented for you.
 2. **Predict** which pattern statistic will tell them apart (spoiler: the **A** statistic)
 3. Run pavo's `classify()` and `adjacent()` on all three demos; compare predictions to measurements
 4. **Your turn** — segment your three lineage fish using `segment_fish()` (or manual crop), then run the same pipeline
 5. Compare your three fish's metrics to the demo fish — does the A statistic correctly classify your fish by pattern orientation, the way it did for the demo?
 6. Submit a worksheet with three metric tables, a comparison answer, and a biological interpretation
-

Need help?

Post in the EEB 187 channel before Wednesday. Don't show up Wednesday with a broken install — you'll spend the lab on installation instead of biology.